

# Artificial Intelligence

## Lecture 03: Solving Problems by Searching

Dr. Bilal Jan

Department of Computer Science / NUCES-FAST, Peshawar

Spring 2026

# Outline

- 1 Problem-Solving Agents
- 2 Search Algorithms
- 3 Uninformed Search
- 4 Informed Search
- 5 Memory-Bounded Search
- 6 Heuristic Functions
- 7 Summary

# Problem-Solving Agents

Artificial Intelligence | Lecture 03

# Representation of Agent's states

Atomic:

- State is a black box with no internal structure
- Simple but uninformative

Factored:

- State is a vector of attribute values (Boolean, real-valued, or symbolic)
- Allows analysis of individual components

Structured:

- State includes objects with attributes and relationships to other objects
- Most expressive representation







Figure: Atomic, Factored and Structured

# Agent's Environment: Goal from Arad → Bucharest



# Problem-Solving Agents

- **Assumed environment:** Single agent, episodic, fully observable, deterministic, static, discrete, and known.
- **Problem-solving agent:** Plans ahead to find action sequences leading to goals
- Uses **atomic representations:** States treated as wholes, no internal structure
- **Four-phase process:**
  - ① **Goal formulation:** Adopt objectives to achieve
  - ② **Problem formulation:** Describe states and actions
  - ③ **Search:** Simulate sequences to find solution
  - ④ **Execution:** Execute actions in solution

# Example: Route Finding in Romania

## Scenario:

- Agent in Arad
- Goal: Reach Bucharest
- Has map of Romania
- Must plan route

## Solution:

- Search through possible routes
- Find: Arad → Sibiu → Fagaras → Bucharest
- Execute actions

## Open-loop vs. Closed-loop

Open-loop: Execute without monitoring percepts (deterministic)

Closed-loop: Monitor percepts during execution (nondeterministic)

# Search Problems: Formal Definition

A search problem consists of:

- ① **State space:** Set of possible states
- ② **Initial state:** Starting state
- ③ **Actions:** Available actions in each state
- ④ **Transition model:** Result of applying action to state
- ⑤ **Goal test:** Check if state is a goal
- ⑥ **Action cost function:** Cost of each action
  
- **Solution:** Sequence of actions from initial to goal state
- **Optimal solution:** Solution with lowest path cost

# Search Problems: Formal Definition

- The **actions** available to the agent. Given a state  $s$ ,  $\text{ACTIONS}(s)$  returns a finite<sup>2</sup> set of [Action](#) actions that can be executed in  $s$ . We say that each of these actions is **applicable** in  $s$ . [Applicable](#)  
An example:

$$\text{ACTIONS}(\text{Arad}) = \{\text{ToSibiu}, \text{ToTimisoara}, \text{ToZerind}\}.$$

- A **transition model**, which describes what each action does.  $\text{RESULT}(s, a)$  returns the [Transition model](#) state that results from doing action  $a$  in state  $s$ . For example,

$$\text{RESULT}(\text{Arad}, \text{ToZerind}) = \text{Zerind}.$$

- An **action cost function**, denoted by  $\text{ACTION-COST}(s, a, s')$  when we are programming [Action cost function](#) or  $c(s, a, s')$  when we are doing math, that gives the numeric cost of applying action  $a$  in state  $s$  to reach state  $s'$ . A problem-solving agent should use a cost function that reflects its own performance measure; for example, for route-finding agents, the cost of an action might be the length in miles (as seen in Figure 3.1), or it might be the time it takes to complete the action.

[Figure: Ref. section 3.1.1 from book](#)

# The Vacuum Cleaner: State space

Agent at LEFT      Agent at RIGHT

-----

(L, D, D) --R--> (R, D, D)

| S                          | S

(L, C, D) --R--> (R, C, D)

(L, D, C) --R--> (R, D, C)


| S                          | S

(L, C, C) --R--> (R, C, C)

Figure: State space

## Extended Vacuum world: 4 squares

What you think about the state space? Assume 4 squares.



## 4 square vacuum world.

- Locations:  $Loc = \{A, B, C, D\}$
- Dirt status per square:  $\{CLEAN, DIRTY\}$
- Any state is 5 tuple:  $s = (loc, a, b, c, d)$
- Here:  $loc \in Loc$  and  $a, b, c, d \in \{CLEAN, DIRTY\}$
- $S = Loc \times \{CLEAN, DIRTY\}^4$
- State space:  $|S| = 4 \times 2^4 = 64$
- n-square world?

# Sliding Puzzles

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

15	2	1	12	8	5	6	11
9	13	7	11	12	1	15	14
11	12	4	15	1	14	9	6
4	15	14		4	9	12	3
13	10	7	8	2	13	7	9
1	3	6	9	5	11	10	15
14	7	15	5	10	2	8	13
8	5	10	3	14	6	4	3

Start State




1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	

Goal State

Figure: Sokoban 3x3:  $9! = 362,880$

Figure: 8x8 Puzzle

# State Space–Graph–Search–Tree



**SEARCH TREE**  
(Each node has ONE unique path to root)

**STATE SPACE**  
(Multiple paths possible)

# Search Algorithms

Artificial Intelligence | Lecture 03

# Best-First Search Framework

- **General approach:** Choose node with minimum evaluation function  $f(n)$
- **Key data structures:**
  - **Node:** Contains STATE, PARENT, ACTION, PATH-COST
  - **Frontier:** Priority queue of nodes to expand
  - **Reached:** Set/table of visited states
- **Algorithm:**
  - ① Choose node with minimum  $f(n)$  from frontier
  - ② If goal, return solution
  - ③ Otherwise, expand node and add children to frontier

# Redundant Paths

- **Problem:** Multiple paths can reach same state
- **Cycle:** Path that returns to previous state
- **Redundant path:** Worse way to reach same state
- **Three approaches:**
  - ① **Graph search:** Remember all reached states
  - ② **Tree-like search:** Don't check for redundant paths
  - ③ **Cycle checking:** Check only for cycles

# Measuring Search Performance

## Four evaluation criteria:

- ① **Completeness:** Guaranteed to find solution?
- ② **Cost optimality:** Finds lowest-cost solution?
- ③ **Time complexity:** How long to find solution?
- ④ **Space complexity:** How much memory needed?

## Complexity measures:

- $b$ : Branching factor;  $d$ : Depth of shallowest solution;  $m$ : Max depth

# Uninformed Search

Artificial Intelligence | Lecture 03

# Uninformed Search Strategies

- **Uninformed search:** No clue about distance to goal
- Only use problem definition (states, actions, costs)
- **Main algorithms:** BFS, UCS, DFS, DLS, IDS, Bidirectional

# Breadth-First Search (BFS)

- **Strategy:** Expand shallowest nodes first (FIFO)
- **Properties:**
  - Complete: Yes; Optimal: Yes (unit costs)
  - Time:  $O(b^d)$ ; Space:  $O(b^d)$
- **Disadvantages:** Exponential space complexity

# Uniform-Cost Search (UCS)

- **Strategy:** Expand node with lowest path cost  $g(n)$
- **Properties:**
  - Complete: Yes; Optimal: Yes
  - Time/Space:  $O(b^{1+\lfloor C^*/\epsilon \rfloor})$
- Uses **late goal test**: Check when node expanded

# Depth-First Search (DFS)

- **Strategy:** Expand deepest unexpanded node first (LIFO)
- **Properties:**
  - Complete: No; Optimal: No
  - Time:  $O(b^m)$ ; Space:  $O(bm)$  (Linear!)

## Depth-Limited Search:

- DFS with depth limit  $l$
- Avoids infinite paths
- Space:  $O(b^l)$

## Iterative Deepening:

- Repeatedly apply DFS with increasing limits
- Complete and optimal (unit costs)
- Space:  $O(bd)$

# Bidirectional Search

- **Idea:** Search forward from start and backward from goal
- **Motivation:**  $b^{d/2} + b^{d/2} \ll b^d$
- **Properties:** Time  $O(b^{d/2})$ , Space  $O(b^{d/2})$

# Comparison of Uninformed Search

Algorithm	Complete?	Optimal?	Time	Space
BFS	Yes	Yes	$O(b^d)$	$O(b^d)$
UCS	Yes	Yes	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
DFS	No	No	$O(b^m)$	$O(bm)$
IDS	Yes	Yes	$O(b^d)$	$O(bd)$

# Informed Search

Artificial Intelligence | Lecture 03

# Informed (Heuristic) Search

- **Informed search:** Uses domain-specific hints (heuristics)
- **Heuristic function  $h(n)$ :** Estimated cost from  $n$  to goal
- Example: Straight-line distance in route-finding

# Greedy Best-First Search

- **Evaluation function:**  $f(n) = h(n)$
- Expands node that appears closest to goal
- **Properties:** Not optimal, can get stuck in loops

- **Evaluation function:**  $f(n) = g(n) + h(n)$
- **Properties** (with admissible  $h$ ):
  - Complete: Yes; Optimal: Yes
  - Optimally efficient: Expands fewest nodes
  - Space:  $O(b^d)$  (Main limitation)

## Admissible Heuristic

$h(n)$  never overestimates the cost to reach goal:  $h(n) \leq h^*(n)$

## Consistent (Monotonic) Heuristic

For every node  $n$  and successor  $n'$ :  $h(n) \leq c(n, a, n') + h(n')$

# Weighted A\* Search

- **Evaluation function:**  $f(n) = g(n) + W \times h(n)$  where  $W > 1$
- **Satisficing search:** Accept "good enough" solutions faster
- Finds solution with cost between  $C^*$  and  $W \times C^*$

# Memory-Bounded Search

Artificial Intelligence | Lecture 03

# Memory-Bounded Search Algorithms

- **Problem:** A\* uses too much memory
- **Main algorithms:** Beam search, IDA\*, RBFS, SMA\*
- **Tradeoff:** Save memory by revisiting states

# Beam Search & IDA\*

- **Beam Search:** Keep only  $k$  best nodes in frontier (Incomplete)
- **IDA\***: Iterative deepening with  $f$ -cost cutoff
  - Complete and optimal; Space:  $O(bd)$

# Heuristic Functions

Artificial Intelligence | Lecture 03

# Designing Heuristic Functions

- **Relaxed problems:** Fewer restrictions (e.g., Manhattan distance)
- **Pattern databases:** Precompute costs for subproblems
- **Landmarks:** Precompute distances to major intermediate states

# Summary

Artificial Intelligence | Lecture 03

# Chapter 3 Summary

- **Uninformed search:** BFS, UCS, DFS, IDS (Blind search)
- **Informed search:** A\*, Greedy, Weighted A\* (Uses  $h(n)$ )
- **Heuristics:** Crucial for performance; must be admissible for optimality
- **Memory:** Main bottleneck for A\*; solved by IDA\*/RBFS